

In [1]:

```
import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")
import seaborn as sns
#Seaborn is a Python data visualization library
```

In [2]:

```
df=pd.read_csv("carIndia.csv")
```

In [3]:

df

Out[3]:

	model_year	maker	model_name	city	distance_covered (km)	fuel_type	pre_owner
0	2012	Maruti	Alto K10 VXI	Mumbai	29067	Petrol	2nd Owner
1	2011	Hyundai	i20 SPORTZ 1.2 O	Mumbai	36791	Petrol	2nd Owner
2	2010	Maruti	A Star VXI	Mumbai	35171	Petrol	1st Owner
3	2011	Hyundai	Santro Xing GLS	Mumbai	19908	Petrol	1st Owner
4	2012	Hyundai	Santro Xing GLS	Mumbai	43847	Petrol	3rd Owner
...
3360	2014	Honda	City S MT DIESEL	Kolkata	61643	Diesel	3rd Owner
3361	2006	Maruti	Wagon R LXI	Kolkata	26500	Petrol	1st Owner
3362	2016	Maruti	S Cross ZETA 1.3	Kolkata	57828	Diesel	1st Owner
3363	2012	BMW	3 Series 320D	Kolkata	23782	Diesel	2nd Owner
3364	2016	Hyundai	Creta 1.4 BASE	Kolkata	33130	Diesel	1st Owner

3365 rows × 8 columns



In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3365 entries, 0 to 3364
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   model_year                            3365 non-null   int64
1   maker                                 3365 non-null   object
2   model_name                            3365 non-null   object
3   city                                  3365 non-null   object
4   distance_covered (km)                 3365 non-null   int64
5   fuel_type                             3365 non-null   object
6   pre_owner                             3365 non-null   object
7   price (₹)                             3365 non-null   int64
dtypes: int64(3), object(5)
memory usage: 210.4+ KB
```

In [5]: `df['model_name'].unique()`

```
'EcoSport AMBIENTE 1.0 TSI', 'Swift Lxi 1.5', 'Duster 85 PS RXL',
'Polo COMFORTLINE 1.2L DIESEL', 'Tiago XZ 1.05 REVOTORQ',
'Amaze 1.2 SAT I VTEC', 'Beat LS DIESEL', 'Etios Liva D 4D G
D',
'Creta 1.6 S', 'Rexton RX7', 'i20 SPORTZ 1.4 CRDI',
'City 1.5 E MT PETROL', 'Etios GD', 'City V CVT',
'Innova Crysta Touring Sport Diesel MT',
'Elite i20 MAGNA 1.4 CRDI', 'Santro Xing GL',
'Creta 1.6 SX (0) CRDI', 'Superb ELEGANCE 1.8 TSI AT',
'Corolla H2 1.8 E', 'Ecosport 1.5 AMBIENTE TDCI',
'VENUE SX(0) CRDi', 'Creta 1.6 SX AT PETROL',
'Duster RXZ AMT 110 PS', 'Vento HIGHLINE PETROL AT',
'S Cross ALPHA 1.3', 'City S MT DIESEL',
'Grand i10 MAGNA 1.1 CRDI', 'Jetta TRENDLINE 1.4 TSI MT',
'Duster 85 PS RXL OPT', 'S Cross ALPHA SHVS',
'Punto EVO MULTIJET 1.3 90 HP', 'i10 ERA', 'Dzire VDI AMT',
'Innova 2.5 GX 7 STR BS IV', 'Etios Liva GD exclusive',
'Benz E Class E 220 CDI ELEGANCE', 'Swift Dzire VXI AT',
'Cruze LTZ', 'Camry HYBRID', 'Vitara Brezza ZDI PLUS',
'Etios Liva D 4D GD SP', 'Jetta HIGHLINE TDI AT',
'Innova ELEGANCE 1.6 EX CRDI', 'Swift DELTA 1.2 DTEC SUNCE'
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	model_year	distance_covered (km)	price (₹)
count	3365.000000	3365.000000	3.365000e+03
mean	2013.876374	60937.813967	4.336549e+05
std	3.035588	41342.775191	2.595909e+05
min	2001.000000	60.000000	2.700000e+04
25%	2012.000000	30598.000000	2.769990e+05
50%	2014.000000	53488.000000	3.647990e+05
75%	2016.000000	82414.000000	4.975990e+05
max	2021.000000	428123.000000	3.600000e+06

```
In [7]: df.isnull().sum()
```

```
Out[7]: model_year      0
maker      0
model_name  0
city      0
distance_covered (km)  0
fuel_type   0
pre_owner   0
price (₹)    0
dtype: int64
```

```
In [8]: df.shape
```

```
Out[8]: (3365, 8)
```

```
In [9]: df.groupby(['model_name']).count()
```

Out[9]:

	model_year	maker	city	distance_covered (km)	fuel_type	pre_owner	price (₹)
model_name							
3 Series 320 D PERFORMANCE EDITION	1	1	1	1	1	1	1
3 Series 320D	2	2	2	2	2	2	2
3 Series 320D LUXURYLINE	2	2	2	2	2	2	2
3 Series 320D SPORTLINE	1	1	1	1	1	1	1
5 Series 520D 2.0	3	3	3	3	3	3	3
...
i20 Magna O 1.4 CRDI	1	1	1	1	1	1	1
i20 SPORTZ 1.2 O	21	21	21	21	21	21	21
i20 SPORTZ 1.2 VTVT	15	15	15	15	15	15	15
i20 SPORTZ 1.4 CRDI	8	8	8	8	8	8	8
i20 SPORTZ O 1.4 CRDI	5	5	5	5	5	5	5

677 rows × 7 columns

```
In [10]: df1=df.drop(['maker','city'],axis=1)
```

```
In [11]: df1
```

Getting the maximum number of cars

SPLITTING

```
In [12]: ▶ def split_model(model):
            parts = model.split(' ')
            Model = ' '.join(parts[:1])
            suffix = parts[-1]
            return Model, suffix
df[['Model', 'Suffix']] = df['model_name'].apply(lambda x: pd.Series(split
#split_model is a custom function defined to split a string (representing
#model.split(' ') splits the input string model into a list of substrings
#Model = ' '.join(parts[:1]): This line attempts to join the first part of
#suffix = parts[-1] retrieves the last part of the split string as the 'Su
#Finally, the function returns a tuple containing 'Model' and 'Suffix'.
#Lambda is a keyword used to create anonymous functions, which are funcio
```

```
In [13]: ▶ df[['Model', 'Suffix']]
```

Out[13]:

	Model	Suffix
0	Alto	VXI
1	i20	O
2	A	VXI
3	Santro	GLS
4	Santro	GLS
...
3360	City	DIESEL
3361	Wagon	LXI
3362	S	1.3
3363	3	320D
3364	Creta	BASE

3365 rows × 2 columns

```
In [ ]: ▶
```

```
In [14]: ▶ df[['Model', 'Suffix']].isnull().sum()
```

Out[14]: Model 0
Suffix 0
dtype: int64

In [15]: ▶

```
# Grouping by 'Model' and 'Suffix' and getting the count of occurrences
result = df[['Model', 'Suffix']].groupby(['Model']).size().reset_index(name='Count')

# Sorting the result by 'Count' in ascending order
sorted_result = result.sort_values('Count', ascending=False)

print(sorted_result)
```

	Model	Count
98	Swift	465
9	Alto	393
110	Wagon	234
121	i10	209
50	Grand	133
..
86	S60	1
90	Sail	1
33	E20	1
93	Scala	1
61	Laura	1

[123 rows x 2 columns]

extracting the maximum and predicting

Swift cars(max)

```
In [16]: ▶ df_Swift=df.loc[(df.Model)=='Swift']
df_Swift.isnull().sum()
```

```
Out[16]: model_year      0
maker      0
model_name  0
city       0
distance_covered (km)  0
fuel_type   0
pre_owner   0
price (₹)    0
Model       0
Suffix      0
dtype: int64
```

```
In [17]: df_Swift=df_Swift.drop(['maker','city','Suffix','model_name'],axis=1)
df_Swift
```

```
Out[17]:
```

	model_year	distance_covered (km)	fuel_type	pre_owner	price (₹)	Model
10	2009	42533	Petrol	2nd Owner	274899	Swift
12	2015	23070	Petrol	2nd Owner	478799	Swift
28	2012	50581	Petrol	1st Owner	406299	Swift
36	2012	47260	Petrol	1st Owner	364799	Swift
37	2012	50525	Diesel	1st Owner	376799	Swift
...
3299	2015	61537	Diesel	1st Owner	450000	Swift
3322	2017	18140	Petrol	1st Owner	457699	Swift
3345	2013	78261	Petrol	1st Owner	327699	Swift
3348	2007	90265	Diesel	3rd Owner	130000	Swift
3354	2017	43860	Petrol	1st Owner	429999	Swift

465 rows × 6 columns

```
In [18]: df_Alto=df.loc[(df.Model)=='Alto']
df_Alto.isnull().sum()
df_Alto=df_Alto.drop(['maker','city','Suffix','model_name'],axis=1)
df_Alto
```

```
Out[18]:
```

	model_year	distance_covered (km)	fuel_type	pre_owner	price (₹)	Model
0	2012	29067	Petrol	2nd Owner	165199	Alto
6	2010	50742	Petrol	2nd Owner	170399	Alto
7	2015	12657	Petrol	1st Owner	282299	Alto
17	2010	34995	Petrol	1st Owner	165999	Alto
34	2015	8322	Petrol	1st Owner	286399	Alto
...
3319	2006	58542	Petrol + CNG	2nd Owner	85000	Alto
3326	2017	11705	Petrol	1st Owner	283199	Alto
3332	2019	1306	Petrol	1st Owner	390999	Alto
3333	2018	3697	Petrol	1st Owner	316999	Alto
3336	2019	2000	Petrol	1st Owner	381299	Alto

393 rows × 6 columns

```
In [19]: df_Wagon=df.loc[(df.Model)=='Wagon']
df_Wagon.isnull().sum()
df_Wagon=df_Wagon.drop(['maker','city','Suffix','model_name'],axis=1)
df_Wagon
```

```
Out[19]:
```

	model_year	distance_covered (km)	fuel_type	pre_owner	price (₹)	Model
8	2013	13688	Petrol	1st Owner	326199	Wagon
14	2011	18514	Petrol	1st Owner	269399	Wagon
16	2012	20712	Petrol	2nd Owner	258399	Wagon
19	2012	39652	Petrol	3rd Owner	288299	Wagon
21	2014	6858	Petrol	1st Owner	358399	Wagon
...
3314	2014	95432	Petrol + CNG	1st Owner	270000	Wagon
3327	2013	22008	Petrol	1st Owner	332599	Wagon
3331	2014	25852	Petrol	1st Owner	345499	Wagon
3357	2001	72000	Petrol	2nd Owner	38000	Wagon
3361	2006	26500	Petrol	1st Owner	100000	Wagon

234 rows × 6 columns

```
In [20]: combined_df = pd.concat([df_Swift,df_Alto,df_Wagon ], axis=0)
```

```
In [21]: combined_df
```

```
Out[21]:
```

	model_year	distance_covered (km)	fuel_type	pre_owner	price (₹)	Model
10	2009	42533	Petrol	2nd Owner	274899	Swift
12	2015	23070	Petrol	2nd Owner	478799	Swift
28	2012	50581	Petrol	1st Owner	406299	Swift
36	2012	47260	Petrol	1st Owner	364799	Swift
37	2012	50525	Diesel	1st Owner	376799	Swift
...
3314	2014	95432	Petrol + CNG	1st Owner	270000	Wagon
3327	2013	22008	Petrol	1st Owner	332599	Wagon
3331	2014	25852	Petrol	1st Owner	345499	Wagon
3357	2001	72000	Petrol	2nd Owner	38000	Wagon
3361	2006	26500	Petrol	1st Owner	100000	Wagon

1092 rows × 6 columns


In []: ▶

In [22]: ▶ `combined_df=pd.get_dummies(combined_df,dtype=int)`In [23]: ▶ `combined_df`

Out[23]:

	model_year	distance_covered (km)	price (₹)	fuel_type_Diesel	fuel_type_Petrol	fuel_type_F +
10	2009	42533	274899	0	1	
12	2015	23070	478799	0	1	
28	2012	50581	406299	0	1	
36	2012	47260	364799	0	1	
37	2012	50525	376799	1	0	
...	
3314	2014	95432	270000	0	0	
3327	2013	22008	332599	0	1	
3331	2014	25852	345499	0	1	
3357	2001	72000	38000	0	1	
3361	2006	26500	100000	0	1	


1092 rows × 14 columns

In [24]: ▶ `cor_mat=combined_df.corr()`

In [25]: ▶ cor_mat


Out[25]:

	model_year	distance_covered (km)	price (₹)	fuel_type_Diesel	fuel_type_Petr
model_year	1.000000	-0.332116	0.589902	-0.003899	0.016400
distance_covered (km)	-0.332116	1.000000	-0.049384	0.454001	-0.477410
price (₹)	0.589902	-0.049384	1.000000	0.343490	-0.253481
fuel_type_Diesel	-0.003899	0.454001	0.343490	1.000000	-0.855815
fuel_type_Petrol	0.016400	-0.477410	-0.253481	-0.855815	1.000000
fuel_type_Petrol + CNG	0.014366	0.076816	-0.100712	-0.125882	-0.374410
fuel_type_Petrol + LPG	-0.112159	0.108164	-0.082885	-0.043086	-0.128116
pre_owner_1st Owner	0.185190	-0.027847	0.155933	0.062116	-0.050410
pre_owner_2nd Owner	-0.159739	0.016274	-0.122632	-0.057498	0.050910
pre_owner_3rd Owner	-0.085008	0.023381	-0.102872	-0.025504	0.021310
pre_owner_4th Owner	0.002483	0.036118	0.004611	0.028439	-0.068310
Model_Alto	0.077723	-0.305099	-0.517942	-0.402212	0.385110
Model_Swift	0.042038	0.336850	0.639359	0.622880	-0.485010
Model_Wagon	-0.141573	-0.049043	-0.164617	-0.280131	0.133910



```
In [26]: ▶ #import seaborn as sns
sns.heatmap(cor_mat,vmax=1,vmin=-1,annot=True,linewidth=0,cmap='bwr')
#cmap is used for vizualization
```

Out[26]: <Axes: >



```
In [27]: ▶ y=combined_df['price (₹)']
print(y)
```

```
10      274899
12      478799
28      406299
36      364799
37      376799
...
3314    270000
3327    332599
3331    345499
3357     38000
3361    100000
Name: price (₹), Length: 1092, dtype: int64
```

```
In [28]: ▶ combined_df.shape
```

Out[28]: (1092, 14)

```
In [29]: ▶ X=combined_df.drop('price (₹)',axis=1)  
print(X)
```

	model_year	distance_covered (km)	fuel_type_Diesel	fuel_type_Pet
rol \				
10	2009	42533	0	
1				
12	2015	23070	0	
1				
28	2012	50581	0	
1				
36	2012	47260	0	
1				
37	2012	50525	1	
0				
...	
...				
3314	2014	95432	0	
0				
3327	2013	22008	0	
1				
3331	2014	25852	0	
1				
3357	2001	72000	0	
1				
3361	2006	26500	0	
1				

	fuel_type_Petrol + CNG	fuel_type_Petrol + LPG	pre_owner_1st Owner
r \			
10	0	0	
0			
12	0	0	
0			
28	0	0	
1			
36	0	0	
1			
37	0	0	
1			
...	
...			
3314	1	0	
1			
3327	0	0	
1			
3331	0	0	
1			
3357	0	0	
0			
3361	0	0	
1			

	pre_owner_2nd Owner	pre_owner_3rd Owner	pre_owner_4th Owner	\
10	1	0	0	
12	1	0	0	
28	0	0	0	
36	0	0	0	
37	0	0	0	
...	

3314	0	0	0
3327	0	0	0
3331	0	0	0
3357	1	0	0
3361	0	0	0

	Model_Alto	Model_Swift	Model_Wagon
10	0	1	0
12	0	1	0
28	0	1	0
36	0	1	0
37	0	1	0
...
3314	0	0	1
3327	0	0	1
3331	0	0	1
3357	0	0	1
3361	0	0	1


[1092 rows x 13 columns]

```
In [30]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_s
```

```
In [31]: X_test.head(5)
```

Out[31]:

	model_year	distance_covered (km)	fuel_type_Diesel	fuel_type_Petrol	fuel_type_Petrol + CNG
217	2011	82414	0	1	0
736	2012	102746	1	0	0
612	2009	74844	0	1	0
3172	2019	61704	0	1	0
2039	2015	53317	0	1	0



```
In [32]: X_train.shape
```

Out[32]: (731, 13)

In [33]: `y_train`

Out[33]:

2510	247299
1908	391799
515	379699
508	399999
2195	352099
...	
2070	393899
6	170399
693	189099
2518	490199
16	258399

Name: price (₹), Length: 731, dtype: int64

In [34]: `y_train.shape`

Out[34]: (731,)

In [35]: `from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train,y_train)
LinearRegression()`

Out[35]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [36]: `ypred=reg.predict(X_test)`

In [37]: ▶ ypred


```
Out[37]: array([195539.36996269, 371699.9915188 , 324974.83253337, 372023.6435042
8,
      288041.70348188, 267115.14946891, 526562.2076845 , 379910.7954368
4,
      250510.56655651, 198464.03027955, 265901.20588164, 241589.9378731
6,
      379730.50353812, 357774.47395833, 195335.10010635, 336269.4230663
1,
      472187.63673207, 196009.44182672, 199329.38984743, 208995.2907243
1,
      348176.17478424, 460394.52768942, 369432.14481308, 371844.0569791
3,
      398324.41459649, 277905.43543871, 310763.46081751, 500630.3430110
1,
      424435.56735088, 300887.34073452, 507442.97542729, 208183.1211062
4,
      218802.41753968, 413497.86382763, 411618.7129089 , 247149.7571234
7,
      299683.53460184, 341757.98877269, 398681.06002911, 287084.0577299
1,
      357405.31299671, 376696.96359129, 268588.27595057, 295892.6764566
2,
      245629.64636811, 411854.5530712 , 237688.99259116, 186357.5660186
9,
      445577.66637183, 373541.6367037 , 208895.04812662, 422545.1592638
7,
      331105.54251634, 203016.07317305, 293560.04711973, 230782.3258485
,
      376518.22812511, 207505.99738307, 271268.05708779, 150969.0283474
8,
      449291.58893806, 338091.49870601, 283910.4957249 , 297590.2093524
9,
      351753.9369653 , 376657.56760898, 178278.7259516 , 180049.7835105
4,
      401189.67283436, 419464.45781165, 489711.68262935, 312558.3524506
5,
      300088.38690739, 284227.2566974 , 231208.43031457, 201357.9867786
8,
      461503.33356599, 289624.09648616, 397252.02539693, 312172.9439679
8,
      359943.66987058, 444532.67637746, 446571.95996616, 249027.2112383
5,
      229846.04634757, 452927.93165959, 159571.37484411, 247009.0165174
2,
      256247.63326975, 169689.79059616, 455444.90736096, 406228.9240977
2,
      402489.92912628, 425914.53480088, 485813.42417437, 246997.6884208
4,
      428618.77865295, 264264.13742673, 446304.53816669, 487288.4223976
2,
      350726.15755045, 441968.52572034, 205931.76284677, 144227.3158642
7,
      352000.27348242, 192916.09392791, 298762.15239481, 379006.4580846
5,
      251177.50335263, 445969.59194408, 443778.12055088, 264959.1380422
6,
      397485.07308071, 246531.32185064, 373321.37575774, 443831.5953272
```

3,
2,
4,
3,
5,
5,
6,
,
4,
6,
4,
3,
3,
1,
2,
3,
1,
6,
5,
7,
,
3,
6,
8,
7,
9,
4,
9,
6,

336000.62331974, 310494.80123378, 250868.29680867, 429342.7893866
433507.91409022, 164150.36167596, 434660.64782907, 394835.5196007
403798.93537169, 157649.26281208, 415673.70852392, 414246.3873519
307465.60732548, 378541.56377341, 200023.91984352, 141087.6942947
306693.92452379, 230922.45225444, 153496.91480777, 285409.1339258
336094.29065692, 321869.48007991, 245966.97066512, 313237.6521938
315473.11558623, 265633.78663628, 318815.14587458, 258877.632126
349133.25853489, 460364.41810273, 249041.19930235, 405527.0523315
253520.44948988, 323073.17817434, 376921.2442832 , 271281.3632293
252884.96875146, 371288.96528591, 293358.09307985, 341218.5881279
380241.527817 , 511530.87310753, 364668.91661638, 222342.0495835
419377.97597047, 398641.2812205 , 270095.04550287, 357604.4275147
311354.637048 , 233736.41382962, 407769.54559939, 98262.6701804
395661.15787908, 316765.02629025, 391699.9017693 , 318812.3137731
169353.37552906, 259544.28768098, 508590.54588293, 270248.1284350
411066.63735847, 427877.56217891, 368476.08956157, 552274.0387195
309469.73165366, 465821.89081709, 346620.47557388, 363847.2205344
429169.12326761, 357265.18355215, 266332.07520799, 505632.4713668
370850.94377624, 443186.8938012 , 465196.39815734, 295499.7784625
354997.91311735, 183208.67552912, 406518.8966507 , 467293.0823323
407446.76577497, 205369.28183891, 349165.64423499, 335635.4618589
210408.23400614, 555675.69512435, 217404.1015026 , 227458.7409081
155809.54108662, 404096.53837916, 385208.31851414, 417274.0754752
312757.80610828, 232182.06864027, 393613.34481183, 296837.3516887
439625.49775854, 361494.36083205, 212130.48845214, 369854.2547992
479903.13765605, 337553.81531535, 394906.26974539, 384890.4999185
320188.94217157, 273821.17467812, 80875.0839765 , 381430.3088568
326803.64349989, 288852.73656257, 285873.82997659, 444568.5261701

295692.29691274, 198324.91810316, 315082.21509865, 352864.8936087
9,
463795.04755688, 226395.29152752, 318267.73936945, 204979.7128631
,
410086.49072031, 291598.03462312, 309125.65992098, 385752.5605467
1,
251280.615917 , 290017.91469362, 463782.38312691, 276117.7677775
5,
349560.29749696, 483251.6041111 , 416321.8844509 , 250930.4034679
5,
451104.36607328, 400460.01583071, 234516.93339521, 357306.0846285
5,
199584.53448899, 439886.67405051, 441647.34025425, 223161.8426562
3,
295884.16108032, 381151.89105947, 217148.52833141, 346657.3208667
9,
414458.97736271, 315515.06246999, 184585.47692198, 440667.8248519
7,
555073.4526161 , 381315.39313077, 239745.11038357, 343183.5396173
1,
229281.51957241, 390880.42700659, 297670.44824027, 407786.4349354
1,
453540.74941368, 398651.39822487, 377699.31777961, 334620.3066632
,
437172.1106319 , 210303.21795141, 181216.09883086, 504607.9965645
3,
199232.43455404, 328852.92541163, 403314.28196623, 417339.1328944
9,
375881.08526492, 403197.77888165, 188979.68008263, 362402.5598572
6,
294628.95253173, 335824.58233231, 289447.37292442, 307905.8562066
,
304212.45930815, 245918.7448544 , 423369.21745346, 209751.1133410
1,
464508.50821053, 313268.23305184, 437642.36232284, 255194.9729920
7,
338515.44760229, 337044.41925766, 231297.50959276, 293494.7894233
8,
362754.76034607, 466011.06815767, 504203.3892541 , 363063.0468049
9,
356859.94483848, 382768.01336357, 273565.90838058, 341887.0994200
6,
202561.12525852, 226021.59803407, 184738.12254742, 406294.9987851
,
179836.69140424, 241548.18232332, 279529.38908598, 470977.1194133
2,
247048.39464473, 349530.74752485, 290662.6784132 , 405964.8433277
5,
419839.63745783, 476301.22336329, 293833.95263045, 310118.9218101
7,
414281.2229151 , 289571.93742224, 302926.42538758, 463863.0124927
3,
249128.71268129, 363590.1357771 , 388101.81719233, 397753.4656928
9,
368188.89521752, 205614.05536704, 489256.78284726, 419420.9946028
3,
206308.0617221 , 246852.45197237, 460241.26493034, 347139.7576852

```
4,
    520785.04039565, 363882.45319352, 424699.80363967, 353747.9492970
9,
    247791.56778919, 328276.38226609, 293096.68948041, 437324.5325352
5,
    259243.61420097, 445686.65165331, 268774.9486716 , 391885.3533658
2,
    348443.30855119, 362335.85311476, 442031.26258421, 439827.9712021
7,
    227816.06826322])
```

```
In [38]: ► from sklearn.metrics import r2_score
          #implements several loss, score, and utility functions to measure classifi
          r2_score(y_test,ypred)
          #r square score is used to evauate performance of regression
```

Out[38]: 0.7731910238316071

```
In [39]: ► from sklearn.metrics import mean_squared_error
          #measures how close the points are to the line
          mean_squared_error(ypred,y_test)
```

Out[39]: 2568269005.9400477


```
In [40]: ▶ results=pd.DataFrame(columns=['price (₹)', 'predicted'])
results['price (₹)']=y_test
results['predicted']=ypred
results=results.reset_index()
#it resets indexes from 0
results['ID']=results.index
results.head(15)
```

Out[40]:

	index	price (₹)	predicted	ID
0	217	205099	195539.369963	0
1	736	346399	371699.991519	1
2	612	208699	324974.832533	2
3	3172	302699	372023.643504	3
4	2039	288699	288041.703482	4
5	2787	231099	267115.149469	5
6	1728	641499	526562.207685	6
7	3096	326099	379910.795437	7
8	3312	216000	250510.566557	8
9	3008	162299	198464.030280	9
10	1135	314399	265901.205882	10
11	2655	270599	241589.937873	11
12	3065	316999	379730.503538	12
13	1255	372999	357774.473958	13
14	1106	234999	195335.100106	14

```
In [41]: ▶ import seaborn as sns
# library which is used for making statistical graphs
import matplotlib.pyplot as plt
#library for visualization
sns.lineplot(x='ID',y='price (₹)',data=results.head(50))
sns.lineplot(x='ID',y='predicted',data=results.head(50))
plt.plot()
```

Out[41]: []



Random Forest Regression

```
In [44]: ▶ from sklearn.model_selection import GridSearchCV #GridSearchCV is for para
from sklearn.ensemble import RandomForestRegressor
reg=RandomForestRegressor()
n_estimators=[25,50,75,100,125,150,175,200] #number of decision trees in t
criterion=['mse'] #criteria for choosing nodes default = 'gini'
max_depth=[3,5,10] #maximum number of nodes in a tree default = None (it v
#parameters={'n_estimators': n_estimators, 'criterion':criterion, 'max_depth'
parameters = {'n_estimators': n_estimators, 'criterion': ['friedman_mse',
RFC_reg = GridSearchCV(reg, parameters)
RFC_reg.fit(X_train,y_train)
```

```
Out[44]: GridSearchCV(estimator=RandomForestRegressor(),
                      param_grid={'criterion': ['friedman_mse', 'squared_error',
                                                'absolute_error', 'poisson'],
                                'max_depth': [3, 5, 10],
                                'n_estimators': [25, 50, 75, 100, 125, 150, 17
5, 200]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [ ]: ▶ #GridSearchCV is tool for a machine Learning model to find the best ones o
#model (Let's say it's a Random Forest) that has some settings you can ad
# Random Forest has hyperparameters like the number of trees in the forest
#Now, trying every possible combination of these settings manually to find
# from sklearn.ensemble module importing random forest regressor algorithm
#It belongs to the ensemble Learning methods and is based on constructing
#criterion parameter in scikit-Learn's RandomForestRegressor specifies the
#the default value for criterion is already set to 'mse', which stands for
```